
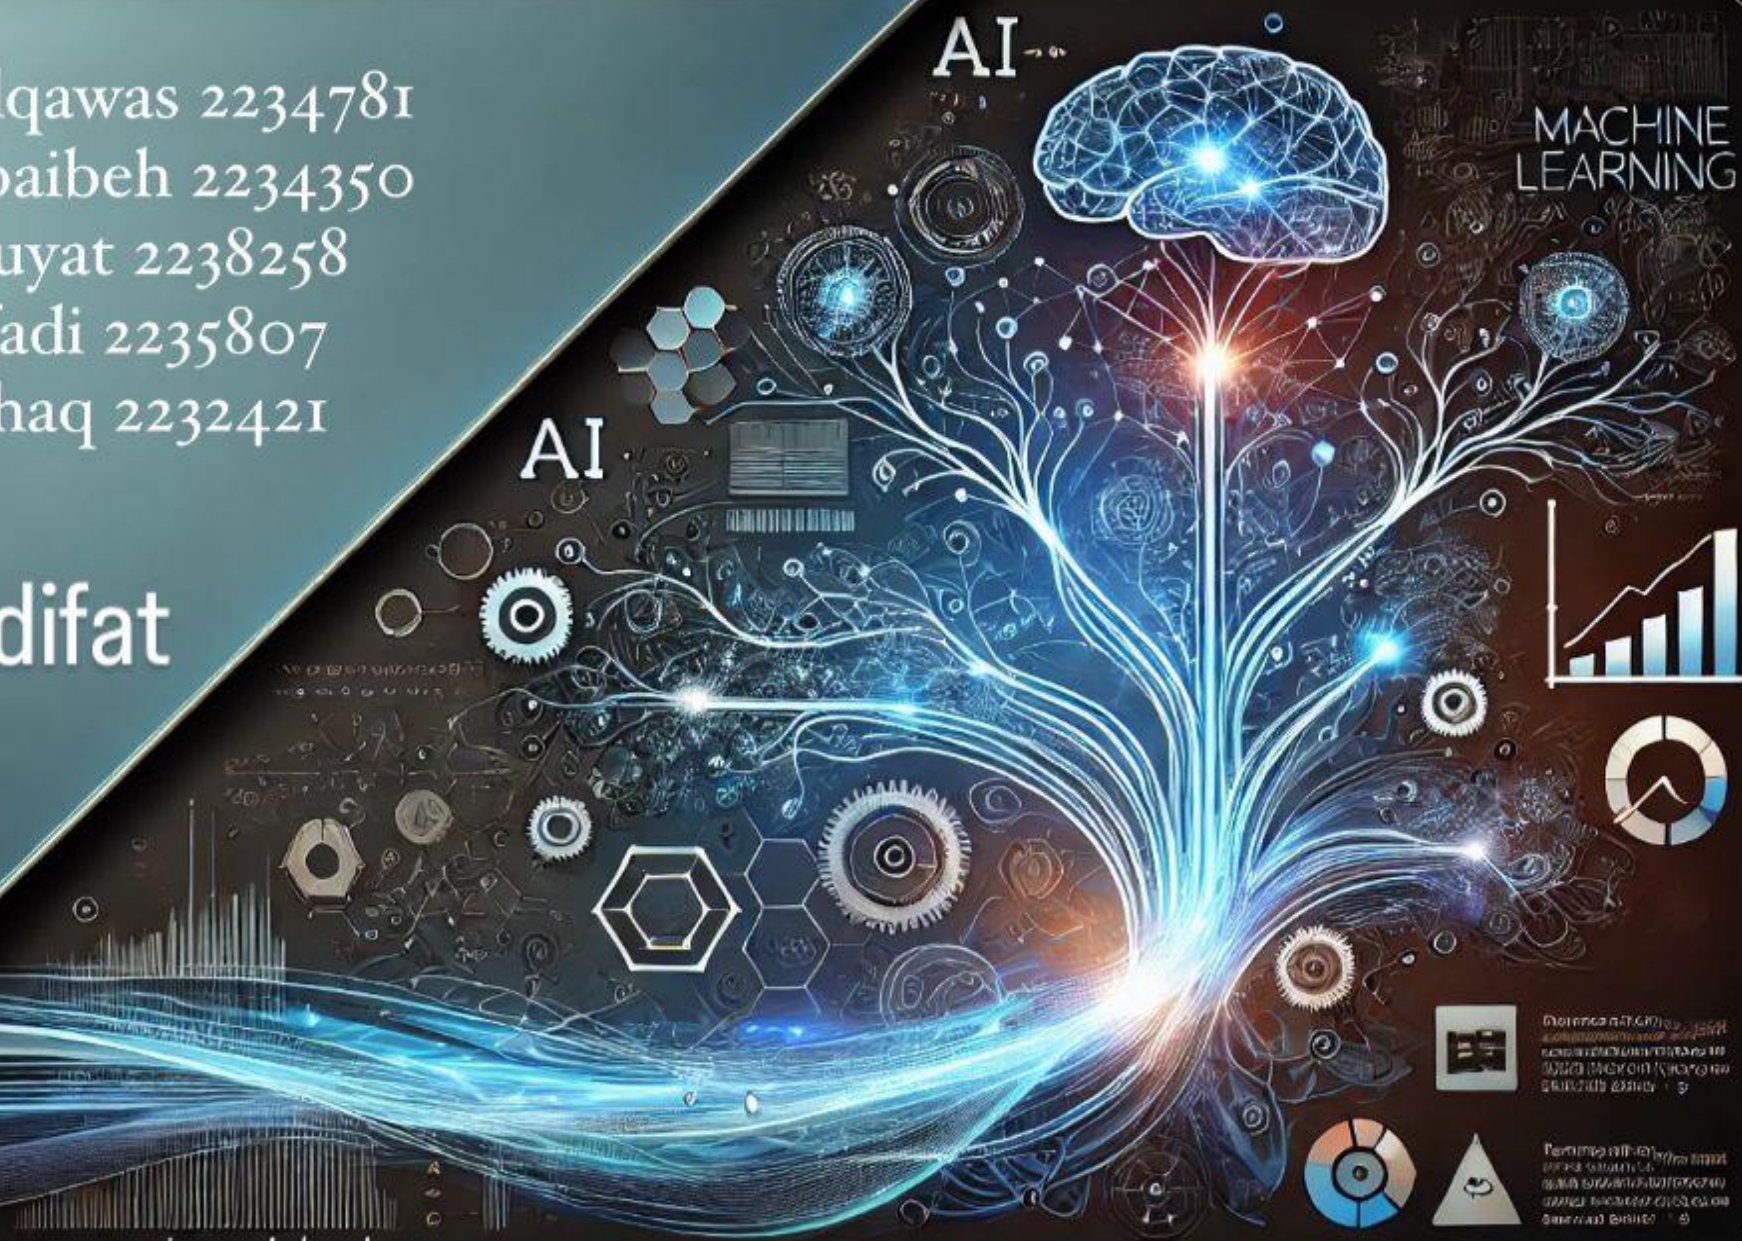




MACHINE LEARNING PROJECT

Amjad Mohammad Alqawas 2234781
Ahmada Shaher Alhabaibeh 2234350
Faten Ayman Alshouyat 2238258
Ahmad Rami Alsafadi 2235807
Mohamad Jamal Ishaq 2232421

Dr.Esra'a Alsdifat



Introduction to the Tasks Data Exploration and Visualization

We will explore the "CO2 Emissions4".

This dataset provides detailed information about a variety of vehicles , including performance specifications, fuel consumption, and CO2 emissions. The columns include data such as vehicle type, engine size, number of cylinders , transmission type, and fuel type. It also records fuel consumption rates in the city, on highways, combined averages, as well as CO2 emissions and fuel efficiency levels.

The dataset contains 7,385 records, covering various types and sizes of vehicles

. All columns are free of missing values, which simplifies handling and analysis

. Additionally, there are 5,692 duplicate rows, indicating the potential for data cleaning to improve analysis quality.

This dataset serves as a solid foundation for analyzing vehicle performance, environmental impact, and operational efficiency, allowing for the exploration of patterns in emissions and fuel consumption across different vehicle categories.

First 5 rows of the dataset:

	Make	Model	'Vehicle Class'	'Engine Size(L)'	Cylinders	Transmission	'Fuel Type'
0	PORSCHE	'911 GT3RS'	TWO-SEATER	4.0	6	AM7	Z
1	BMW	M3	COMPACT	3.0	6	M6	Z
2	VOLKSWAGEN	'GOLF SPORTWAGON TDI (modified)'	'STATION WAGON - SMALL'	2.0	4	M6	D
3	BMW	'335i xDRIVE SEDAN'	COMPACT	3.0	6	AS8	Z
4	MASERATI	'QUATTROPORTE GTS'	FULL-SIZE	3.8	8	AS8	Z

'Fuel Type'	'Fuel Consumption City (L/100 km)'	'Fuel Consumption Hwy (L/100 km)'	'Fuel Consumption Comb (L/100 km)'	'Fuel Consumption Comb (mpg)'	'CO2 Emissions(g/km)'
Z	16.3	11.8	14.3	20	334
Z	13.7	9.0	11.6	24	271
D	7.6	5.5	6.7	42	179
Z	11.9	7.8	10.0	28	230
Z	17.6	10.7	14.5	19	334

Dataset Information:

	Column	Non-Null Count	Data Type
Make	Make	7385	object
Model	Model	7385	object
'Vehicle Class'	'Vehicle Class'	7385	object
'Engine Size(L)'	'Engine Size(L)'	7385	float64
Cylinders	Cylinders	7385	int64
Transmission	Transmission	7385	object
'Fuel Type'	'Fuel Type'	7385	object
'Fuel Consumption City (L/100 km)'	'Fuel Consumption City (L/100 km)'	7385	float64
'Fuel Consumption Hwy (L/100 km)'	'Fuel Consumption Hwy (L/100 km)'	7385	float64
'Fuel Consumption Comb (L/100 km)'	'Fuel Consumption Comb (L/100 km)'	7385	float64
'Fuel Consumption Comb (mpg)'	'Fuel Consumption Comb (mpg)'	7385	int64
'CO2 Emissions(g/km)'	'CO2 Emissions(g/km)'	7385	int64

Missing values per column:

	Column	Missing Values
0	Make	0
1	Model	0
2	'Vehicle Class'	0
3	'Engine Size(L)'	0
4	Cylinders	0
5	Transmission	0
6	'Fuel Type'	0
7	'Fuel Consumption City (L/100 km)'	0
8	'Fuel Consumption Hwy (L/100 km)'	0
9	'Fuel Consumption Comb (L/100 km)'	0
10	'Fuel Consumption Comb (mpg)'	0
11	'CO2 Emissions(g/km)'	0

Column Names :

Column Name	
0	Make
1	Model
2	'Vehicle Class'
3	'Engine Size(L)'
4	Cylinders
5	Transmission
6	'Fuel Type'
7	'Fuel Consumption City (L/100 km)'
8	'Fuel Consumption Hwy (L/100 km)'
9	'Fuel Consumption Comb (L/100 km)'
10	'Fuel Consumption Comb (mpg)'
11	'CO2 Emissions(g/km)'

Number of duplicated rows:

	Metric	Count
0	Duplicated Rows	5692

Dataset Dimensions:


	Metric	Count
0	Rows	7385
1	Columns	12

Part 1: Data Analysis and Visualization


Tasks : 1. Visualization Tasks

Visualize the relationships between Engine Size, Fuel Consumption, and CO₂ Emissions using scatterplots.


Fuel Consumption City (L/100 km) vs CO2 Emissions




Fuel Consumption Hwy (L/100 km) vs CO2 Emissions




Fuel Consumption Comb (L/100 km) vs CO2 Emissions




Fuel Consumption Comb (mpg) vs CO2 Emissions



Engin size vs CO2
Emission :




Engine Size vs Fuel Consumption :

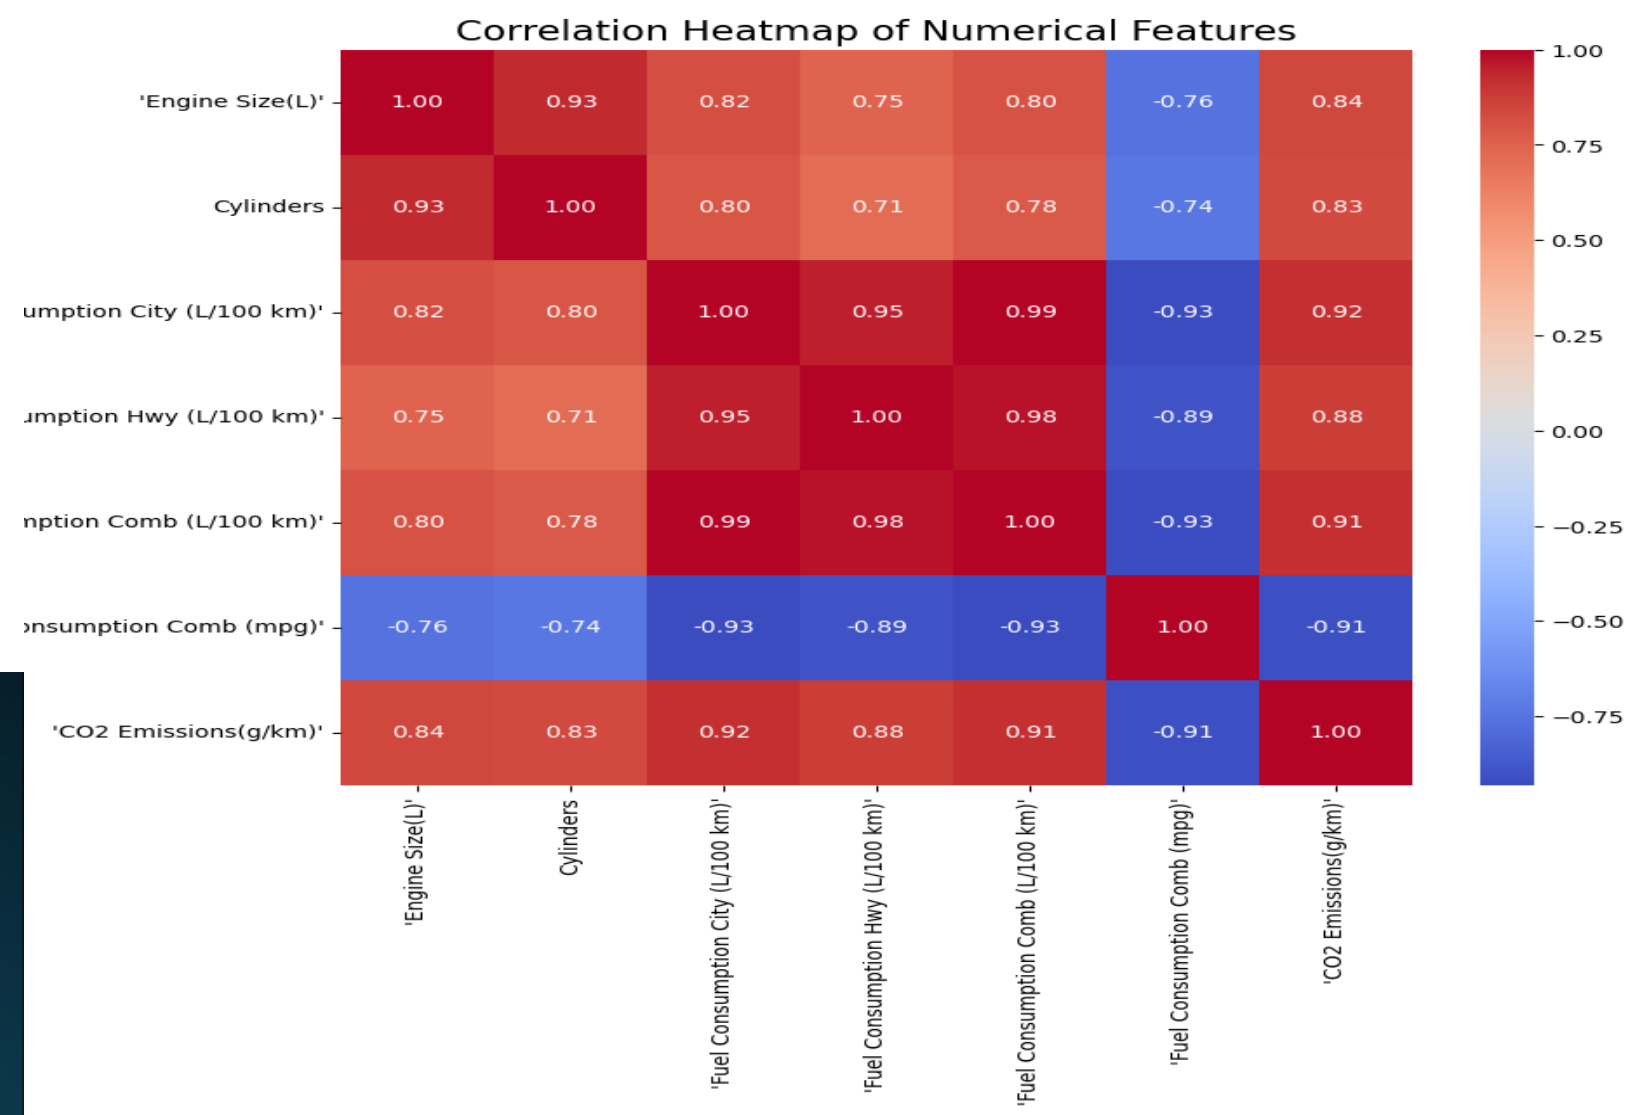


Tasks : 1. Visualization Tasks

Visualize the relationships between Engine Size, Fuel Consumption, and CO₂ Emissions using pair plots.



Create a heatmap of correlations to understand how numerical features are related to CO₂ Emissions.



Key features contributing to CO₂ emissions

2. Feature Insights:

- o Identify key features contributing to CO₂ emissions (e.g., does Fuel Type significantly impact emissions?).

- o Provide a written summary of observed trends or patterns.

Feature Insights: Identifying Key Features Contributing to CO₂ Emissions

1. Key Features Contributing to CO₂ Emissions Based on data analysis and feature correlations, the following are the primary factors contributing to CO₂ emissions:

Fuel Consumption Comb (L/ 100 km):

This is the most impactful feature on emissions. As fuel consumption (L/100 km) increases, CO₂ emissions increase proportionally. Vehicles with lower fuel efficiency, such as sports cars or those with large engines, record higher emissions.

Engine Size (L):

There is a strong positive relationship between engine size and emissions. Larger engines generate more power but also burn more fuel, leading to higher CO₂ emissions. **Cylinders:**

The number of cylinders is positively correlated with emissions. Engines with more cylinders are generally more powerful but less fuel-efficient. **Fuel Consumption Comb (mpg):**

There is a strong negative correlation with emissions. Vehicles with higher efficiency (more miles per gallon) produce significantly lower CO₂ emissions. **Fuel Type:**

Fuel type impacts emissions, with some types (like diesel) being more efficient than gasoline, resulting in lower emissions. Alternative fuels (such as electricity or natural gas) can drastically reduce emissions.

2. Observed Trends and Patterns

Fuel Consumption and CO₂ Emissions:

A clear and strong correlation exists between fuel consumption (city, highway, and combined) and emissions. Higher fuel consumption directly results in higher emissions. **Engine Size and Cylinders:** Vehicles with larger engines and more cylinders emit more CO₂, indicating that these characteristics are associated with higher fuel consumption. **Fuel Type:** Fuel type plays an important role, but its impact depends on the overall efficiency of the vehicle. Electric or alternative-fuel vehicles record significantly lower emissions compared to diesel or gasoline. **High-Efficiency Vehicles:** Vehicles with lower fuel consumption (more miles per gallon) and energy-saving technologies produce the least amount of CO₂ emissions.

Summary :
The primary features affecting CO₂ emissions are fuel consumption, engine size, and the number of cylinders. Vehicles with larger engines and higher cylinder counts consume more fuel and thus produce more emissions. Fuel type has a moderate impact, with alternative fuels outperforming conventional fuels in reducing emissions. The shift towards more efficient technologies, such as electric or hybrid engines, is essential for reducing emissions on a global scale.



Part 2: Regression Analysis

Dataset: Use the same dataset as Part 1.

Tasks:

1 . Model Construction:

- o Split the dataset into training (70%) and testing (30%) subsets.
- o Build a Multiple Linear Regression model to predict CO₂ Emissions using:

- ☒ Engine Size
- ☒ Fuel Consumption (Comb)
- ☒ Cylinders

Intercept: 24.017

Coefficients: 2.968, 18.272, 3.526

Model Performance:

Mean Squared Error (MSE): 176.402

R-squared (R²): 0.929

Predicted CO₂ Emissions for the new record: 210.875

Part 2: Regression Analysis

2 . Feature Selection:

- o Apply feature selection techniques (Forward technique) to refine the regression model.
- o Rebuild the model using the selected features.

Selected Features: ['Fuel_Consumption_Comb_(L/100_km)', 'Cylinders', 'Engine_Size(L)']

Final Model Summary:

OLS Regression Results				
=====				
Dep. Variable:	CO2_Emissions(g/km)	R-squared:	0.878	
Model:	OLS	Adj. R-squared:	0.878	
Method:	Least Squares	F-statistic:	1.237e+04	
Date:	Sun, 05 Jan 2025	Prob (F-statistic):	0.00	
Time:	00:26:12	Log-Likelihood:	-22830.	
No. Observations:	5169	AIC:	4.567e+04	
Df Residuals:	5165	BIC:	4.569e+04	
Df Model:	3			
Covariance Type:	nonrobust			
=====				
=====				
		coef	std err	t P> t
[0.025 0.975]				

const		52.5800	1.295	40.604	0.000
50.041	55.119				
Fuel_Consumption_Comb_(L/100_km)		12.9179	0.166	77.987	0.000
12.593	13.243				
Cylinders		7.5437	0.422	17.879	0.000
6.717	8.371				
Engine_Size(L)		4.5441	0.592	7.679	0.000
3.384	5.704				
=====					
Omnibus:	1311.366	Durbin-Watson:			1.954
Prob(Omnibus):	0.000	Jarque-Bera (JB):			4606.632
Skew:	-1.246	Prob(JB):			0.00
Kurtosis:	6.896	Cond. No.			64.4
=====					

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

=====

Part 2: Regression Analysis

3 . Evaluation Metrics :

Evaluate model performance using:

R² score.

Mean Squared Error (MSE).

Model Evaluation on Test Data:
Mean Squared Error (MSE): 427.475
R-squared (R²): 0.866

4 . Insights:

- o Compare the initial and refined regression models based on evaluation metrics.
- o Interpret the impact of key features on CO₂ emissions (e.g., how does an increase in engine size affect emissions?).

Comparison of Models:

Initial Model - R² Score: 0.866 | MSE: 427.475

Refined Model - R² Score: 0.866 | MSE: 427.475

Impact of Key Features on CO₂ Emissions:

An increase in Fuel_Consumption_Comb_(L/100_km) increases CO₂ emissions by 12.918 units.

An increase in Cylinders increases CO₂ emissions by 7.544 units.

An increase in Engine_Size(L) increases CO₂ emissions by 4.544 units.

Built a **Random Forest Regressor** model to predict target values using training data, then evaluated its performance using **Mean Squared Error (MSE)** and **R-squared (R^2)** metrics.

```
from sklearn.ensemble import RandomForestRegressor

rf = RandomForestRegressor(n_estimators=100, max_depth=10, random_state=42)
rf.fit(X_train, y_train)

y_pred_rf = rf.predict(X_test)
mse_rf = mean_squared_error(y_test, y_pred_rf)
r2_rf = r2_score(y_test, y_pred_rf)

print("\nRandom Forest Model Performance:")
print(f"Mean Squared Error (MSE): {mse_rf:.3f}")
print(f"R-squared ( $R^2$ ): {r2_rf:.3f}")
```

```
Random Forest Model Performance:
Mean Squared Error (MSE): 43.781
R-squared ( $R^2$ ): 0.986
```


Part 3: Classification

A. Initial Data Exploration

First 5 rows of the dataset:

	loan_id	no_of_dependents'	education'	self_employed'	income_annum'	loan_amount'	loan_term'
0	1103	1	'Graduate'	'Yes'	1600000	4900000	16
1	2748	1	'Graduate'	'Yes'	9700000	25500000	18
2	3889	1	'Not Graduate'	'No'	9100000	23300000	18
3	2921	1	'Not Graduate'	'No'	400000	1400000	4
4	3589	4	'Not Graduate'	'Yes'	6900000	22000000	10

cibil_score'	residential_assets_value'	commercial_assets_value'	luxury_assets_value'	bank_asset_value'	loan_status
487	4300000	1100000	3900000	2000000	'Rejected'
741	21900000	6200000	34800000	4900000	'Approved'
792	800000	12600000	32100000	12400000	'Approved'
302	200000	100000	1500000	500000	'Approved'
512	6100000	9100000	20600000	3900000	'Rejected'

Dataset Information:

	Column	Non-Null Count	Data Type
loan_id	loan_id	4269	int64
'no_of_dependents'	'no_of_dependents'	4269	int64
'education'	'education'	4269	object
'self_employed'	'self_employed'	4269	object
'income_annum'	'income_annum'	4269	int64
'loan_amount'	'loan_amount'	4269	int64
'loan_term'	'loan_term'	4269	int64
'cibil_score'	'cibil_score'	4269	int64
'residential_assets_value'	'residential_assets_value'	4269	int64
'commercial_assets_value'	'commercial_assets_value'	4269	int64
'luxury_assets_value'	'luxury_assets_value'	4269	int64
'bank_asset_value'	'bank_asset_value'	4269	int64
'loan_status'	'loan_status'	4269	object

Missing values per column:

	Column	Missing Values
0	loan_id	0
1	'no_of_dependents'	0
2	'education'	0
3	'self_employed'	0
4	'income_annum'	0
5	'loan_amount'	0
6	'loan_term'	0
7	'cibil_score'	0
8	'residential_assets_value'	0
9	'commercial_assets_value'	0
10	'luxury_assets_value'	0
11	'bank_asset_value'	0
12	'loan_status'	0

Column Names :

Column Name	
0	loan_id
1	' no_of_dependents'
2	' education'
3	' self_employed'
4	' income_annum'
5	' loan_amount'
6	' loan_term'
7	' cibil_score'
8	' residential_assets_value'
9	' commercial_assets_value'
10	' luxury_assets_value'
11	' bank_asset_value'
12	' loan_status'

Dataset Dimensions:

	Metric	Count
0	Rows	4269
1	Columns	13

Data types for columns :

loan_id	int64
' no_of_dependents'	int64
' education'	object
' self_employed'	object
' income_annum'	int64
' loan_amount'	int64
' loan_term'	int64
' cibil_score'	int64
' residential_assets_value'	int64
' commercial_assets_value'	int64
' luxury_assets_value'	int64
' bank_asset_value'	int64
' loan_status'	object
dtype: object	

Tasks:

1 . Preprocess the dataset (handle missing values, encode categorical variables).

I handled the missing data, represented by zeros, by calculating the mean value and imputing it into these entries.

Imputed means:

no_of_dependents: 2.98

residential_assets_value: 7792397.24

commercial_assets_value: 5323549.24

bank_asset_value: 5123744.72

Sample of processed data:

	loan_id	no_of_dependents	income_annum	loan_amount	loan_term	cibil_score
0	1103	1.0	1600000	4900000	16	487
1	2748	1.0	9700000	25500000	18	741
2	3889	1.0	9100000	23300000	18	792
3	2921	1.0	400000	1400000	4	302
4	3589	4.0	6900000	22000000	10	512

residential_assets_value	commercial_assets_value	luxury_assets_value	bank_asset_value	education_' Not Graduate'	self_employed_' Yes'	loan_status_' Rejected'
4300000.0	1100000.0	3900000	2000000.0	0	1	1
21900000.0	6200000.0	34800000	4900000.0	0	1	0
800000.0	12600000.0	32100000	12400000.0	1	0	0
200000.0	100000.0	1500000	500000.0	1	0	0
6100000.0	9100000.0	20600000	3900000.0	1	1	1

2 . Apply various classification

First :

DecisionTree

Decision Tree Classifier:
Accuracy: 0.99

	precision	recall	f1-score	support
0	0.99	0.99	0.99	794
1	0.99	0.99	0.99	487
accuracy			0.99	1281
macro avg	0.99	0.99	0.99	1281
weighted avg	0.99	0.99	0.99	1281



second : LogisticRegression

Logistic Regression:

Accuracy: 0.81

	precision	recall	f1-score	support
0	0.81	0.90	0.85	794
1	0.80	0.67	0.73	487
accuracy			0.81	1281
macro avg	0.81	0.78	0.79	1281
weighted avg	0.81	0.81	0.81	1281



Third : Naïve Bayes

Naïve Bayes:

Accuracy: 0.75

	precision	recall	f1-score	support
0	0.73	0.96	0.83	794
1	0.86	0.42	0.57	487
accuracy			0.75	1281
macro avg	0.80	0.69	0.70	1281
weighted avg	0.78	0.75	0.73	1281

Fourth : RandomForestClassifier (Ensamble)

Random Forest Classifier:

Accuracy: 0.99

	precision	recall	f1-score	support
0	0.99	0.99	0.99	794
1	0.99	0.99	0.99	487
accuracy			0.99	1281
macro avg	0.99	0.99	0.99	1281
weighted avg	0.99	0.99	0.99	1281

last classifier : SVM with StandardScaler

Support Vector Machine (with Feature Scaling):

Accuracy: 0.90

	precision	recall	f1-score	support
0	0.95	0.89	0.92	794
1	0.84	0.93	0.88	487
accuracy			0.90	1281
macro avg	0.89	0.91	0.90	1281
weighted avg	0.91	0.90	0.90	1281

3- Evaluate models using confusion matrices and performance metrics (Accuracy, Precision, Recall, F1-Score, AUC).

```
=== Performance Metrics for Decision Tree ===
```

```
Accuracy: 0.99
```


```
Precision: 0.99
```

```
Recall: 0.99
```

```
F1-Score: 0.99
```

```
Classification Report:
```

	precision	recall	f1-score	support
0	0.99	0.99	0.99	794
1	0.99	0.99	0.99	487
accuracy			0.99	1281
macro avg	0.99	0.99	0.99	1281
weighted avg	0.99	0.99	0.99	1281



=== Performance Metrics for Logistic Regression ===

Accuracy: 0.81


Precision: 0.80

Recall: 0.67

F1-Score: 0.73

Classification Report:

	precision	recall	f1-score	support
0	0.81	0.90	0.85	794
1	0.80	0.67	0.73	487
accuracy			0.81	1281
macro avg	0.81	0.78	0.79	1281
weighted avg	0.81	0.81	0.81	1281



=== Performance Metrics for SVM ===

Accuracy: 0.62


Precision: 0.00

Recall: 0.00

F1-Score: 0.00

Classification Report:

	precision	recall	f1-score	support
0	0.62	1.00	0.77	794
1	0.00	0.00	0.00	487
accuracy			0.62	1281
macro avg	0.31	0.50	0.38	1281
weighted avg	0.38	0.62	0.47	1281



=== Performance Metrics for Naive Bayes ===

Accuracy: 0.75


Precision: 0.86

Recall: 0.42

F1-Score: 0.57

Classification Report:

	precision	recall	f1-score	support
0	0.73	0.96	0.83	794
1	0.86	0.42	0.57	487
accuracy			0.75	1281
macro avg	0.80	0.69	0.70	1281
weighted avg	0.78	0.75	0.73	1281



=== Performance Metrics for Random Forest ===

Accuracy: 0.99

Precision: 0.99


Recall: 0.99

F1-Score: 0.99

Classification Report:

	precision	recall	f1-score	support
0	0.99	0.99	0.99	794
1	0.99	0.99	0.99	487
accuracy			0.99	1281
macro avg	0.99	0.99	0.99	1281
weighted avg	0.99	0.99	0.99	1281

Confusion Matrix for Random Forest



3 . Compare hold-out validation (70/30 split) with 10-fold cross validation.

```
=== Holdout Validation ===
```

```
Holdout Accuracy: 0.992
```

```
=== 10-Fold Cross Validation ===
```

```
Cross-Validation Accuracy Scores: 0.995 / 1.000 / 0.995 / 0.991 / 1.000 / 0.998  
/ 0.988 / 0.998 / 0.993 / 0.993
```

```
Mean Accuracy: 0.995
```

```
Standard Deviation: 0.004
```

4 . Identify the classifier (Classification model) that the group has selected (the type, it's performance, reasons for selection).

The type : Random Forest(Ensemble)

it's performance : 99%

Reasons for selection : After building various models, it was found that Decision Tree and Random Forest provided the best accuracy. However, given the nature of the data, which may suffer from feature overlap, this could lead to overfitting in the case of the Decision Tree. Therefore, we chose Random Forest as the best classifier since it delivers optimal performance with medium to large datasets. Our dataset, which is considered medium-sized, achieved a very high accuracy of 99%, making Random Forest an excellent choice.



99%

ACCURECE ACCURACY



THANK YOU
MACHINE LEARNING PROJECT

MACHINE
LEARNING

399